

## 1. Business Understanding

Customer churn refers to the phenomenon where customers or subscribers end their relationship with a company or service provider. The primary objective of any business is to reduce customer churn. By identifying potential churners in advance, SyriaTel can take precaution measures to retain these customers.

Such measures may include:

- Improved customer service.
- Addressing the issues that may lead to churn.
- Making targeted offers and,
- Targeted market campaigns among others.

Churn can have great significant financial impact on the business as high churn leads to loss of recurring revenue, damage of brand's reputation among many other effects.

This project aims to predict customer churn by developing an algorithm to predict the churn rate based on customer usage pattern.

## Objectives

- To determine attributes that contribute to customer churn.
- To build a classification model that predicts customer churn.
- To achieve a recall score for the model of at least 70%
- To make valid recommendations to SyriaTel on ways they can reduce customer churn.

## Analytical Questions

1. Which factors generally lead to customer churn?
2. Which is the most appropriate evaluation metric for the model I will build?
3. What recommendations do I have for the stakeholder?

## 2. Data Understanding

The dataset has customer usage pattern and whether the customer has churned or not. I will develop an algorithm to predict the churn score based on usage pattern. The predictors provided are as follows:

1. "state", string. 2-letter code of the US state of customer residence
2. "account\_length", Number of months the customer has been with current telco provider
3. "area\_code", a string 3 digit area code.
4. "international\_plan", (yes/no). The customer has international plan.
5. "voice\_mail\_plan", (yes/no). The customer has voice mail plan.
6. "number\_vmail\_messages", numerical. Number of voice-mail messages.
7. "total\_day\_minutes", numerical. Total minutes of day calls.
8. "total\_day\_calls", numerical. Total minutes of day calls.
9. "total\_day\_charge", numerical. Total charge of day calls.
10. "total\_eve\_minutes", numerical. Total minutes of evening calls.

11. "total\_eve\_calls", numerical. Total number of evening calls.
12. "total\_eve\_charge", numerical. Total charge of evening calls.
13. "total\_night\_minutes", numerical. Total minutes of night calls.
14. "total\_night\_calls", numerical. Total number of night calls.
15. "total\_night\_charge", numerical. Total charge of night calls.
16. "total\_intl\_minutes", numerical. Total minutes of international calls.
17. "total\_intl\_calls", numerical. Total number of international calls.
18. "total\_intl\_charge", numerical. Total charge of international calls
19. "number\_customer\_service\_calls", numerical. Number of calls to customer service

Target Variable is:

- Churn: if the customer has churned (1=yes; 0 = no)

## Data Exploration and Cleaning

In [1]:

```
# Importing libraries.
import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns

import warnings
warnings.filterwarnings('ignore')

from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.metrics import accuracy_score, recall_score, precision_score, f1_score
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from imblearn.over_sampling import SMOTE
from sklearn.preprocessing import StandardScaler
```

## 2.1 Loading the dataset

In [2]:

```
# Creating a dataframe and viewing the first 5 columns.
df = pd.read_csv('bigml_59c28831336c6604c800002a.csv')
df.head()
```

Out[2]:

|   | state | account length | area code | phone number | international plan | voice mail plan | number vmail messages | total day minutes | total day calls | total day charge | ... |
|---|-------|----------------|-----------|--------------|--------------------|-----------------|-----------------------|-------------------|-----------------|------------------|-----|
| 0 | KS    | 128            | 415       | 382-4657     | no                 | yes             | 25                    | 265.1             | 110             | 45.07            | ... |
| 1 | OH    | 107            | 415       | 371-7191     | no                 | yes             | 26                    | 161.6             | 123             | 27.47            | ... |
| 2 | NJ    | 137            | 415       | 358-1921     | no                 | no              | 0                     | 243.4             | 114             | 41.38            | ... |
| 3 | OH    | 84             | 408       | 375-9999     | yes                | no              | 0                     | 299.4             | 71              | 50.90            | ... |
| 4 | OK    | 75             | 415       | 330-6626     | yes                | no              | 0                     | 166.7             | 113             | 28.34            | ... |

5 rows × 21 columns



## 2.2 Statistical Analysis.

In [3]:

```
# A function to analyze the shape, number of columns, and information of the data
def analyze_dataset(filename):
    """
    This function outputs information about the shape,
    columns, and information of the dataset using the Pandas library.
    """
    # Output the shape of the dataset
    print("Shape of dataset:", df.shape)
    print('\n-----')

    # Output the column names of the dataset
    print("Column names:", list(df.columns))
    print('\n-----')

    # Output information about the dataset
    print(df.info())
    print('\n-----')

    # output descriptive statistics about the dataset
    print(df.describe())
    print('\n-----')

    # output if the dataset has duplicates
    print("Number of duplicates: ",df.duplicated().sum())
```

In [4]:

```
# Applying the function to analyze our dataset  
analyze_dataset('bigml_59c28831336c6604c800002a.csv')
```

Shape of dataset: (3333, 21)

-----  
Column names: ['state', 'account length', 'area code', 'phone number', 'international plan', 'voice mail plan', 'number vmail messages', 'total day minutes', 'total day calls', 'total day charge', 'total eve minutes', 'total eve calls', 'total eve charge', 'total night minutes', 'total night calls', 'total night charge', 'total intl minutes', 'total intl calls', 'total intl charge', 'customer service calls', 'churn']

-----  
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 3333 entries, 0 to 3332  
Data columns (total 21 columns):  
#      Column                                      Non-Null Count     Dtype  
---      -  
0      state                                            3333 non-null     object  
1      account length                                3333 non-null     int64  
2      area code                                      3333 non-null     int64  
3      phone number                                  3333 non-null     object  
4      international plan                            3333 non-null     object  
5      voice mail plan                               3333 non-null     object  
6      number vmail messages                       3333 non-null     int64  
7      total day minutes                            3333 non-null     float64  
8      total day calls                               3333 non-null     int64  
9      total day charge                              3333 non-null     float64  
10     total eve minutes                            3333 non-null     float64  
11     total eve calls                               3333 non-null     int64  
12     total eve charge                              3333 non-null     float64  
13     total night minutes                          3333 non-null     float64  
14     total night calls                            3333 non-null     int64  
15     total night charge                           3333 non-null     float64  
16     total intl minutes                           3333 non-null     float64  
17     total intl calls                              3333 non-null     int64  
18     total intl charge                            3333 non-null     float64  
19     customer service calls                      3333 non-null     int64  
20     churn                                          3333 non-null     bool  
dtypes: bool(1), float64(8), int64(8), object(4)  
memory usage: 524.2+ KB  
None

-----  
                 account length                   area code    number vmail messages    total day  
minutes \  
count       3333.000000    3333.000000                   3333.000000            333  
3.000000  
mean           101.064806       437.182418                   8.099010               17  
9.775098  
std            39.822106        42.371290                   13.688365               5  
4.467389  
min            1.000000        408.000000                   0.000000  
0.000000  
25%            74.000000        408.000000                   0.000000               14  
3.700000  
50%            101.000000       415.000000                   0.000000               17  
9.400000  
75%            127.000000       510.000000                   20.000000               21  
6.400000  
max            243.000000       510.000000                   51.000000               35  
0.800000

|            | total day calls | total day charge | total eve minutes | total e |
|------------|-----------------|------------------|-------------------|---------|
| ve calls \ |                 |                  |                   |         |
| count      | 3333.000000     | 3333.000000      | 3333.000000       | 333     |
| 3.000000   |                 |                  |                   |         |
| mean       | 100.435644      | 30.562307        | 200.980348        | 10      |
| 0.114311   |                 |                  |                   |         |
| std        | 20.069084       | 9.259435         | 50.713844         | 1       |
| 9.922625   |                 |                  |                   |         |
| min        | 0.000000        | 0.000000         | 0.000000          |         |
| 0.000000   |                 |                  |                   |         |
| 25%        | 87.000000       | 24.430000        | 166.600000        | 8       |
| 7.000000   |                 |                  |                   |         |
| 50%        | 101.000000      | 30.500000        | 201.400000        | 10      |
| 0.000000   |                 |                  |                   |         |
| 75%        | 114.000000      | 36.790000        | 235.300000        | 11      |
| 4.000000   |                 |                  |                   |         |
| max        | 165.000000      | 59.640000        | 363.700000        | 17      |
| 0.000000   |                 |                  |                   |         |

|       | total eve charge | total night minutes | total night calls \ |
|-------|------------------|---------------------|---------------------|
| count | 3333.000000      | 3333.000000         | 3333.000000         |
| mean  | 17.083540        | 200.872037          | 100.107711          |
| std   | 4.310668         | 50.573847           | 19.568609           |
| min   | 0.000000         | 23.200000           | 33.000000           |
| 25%   | 14.160000        | 167.000000          | 87.000000           |
| 50%   | 17.120000        | 201.200000          | 100.000000          |
| 75%   | 20.000000        | 235.300000          | 113.000000          |
| max   | 30.910000        | 395.000000          | 175.000000          |

|       | total night charge | total intl minutes | total intl calls \ |
|-------|--------------------|--------------------|--------------------|
| count | 3333.000000        | 3333.000000        | 3333.000000        |
| mean  | 9.039325           | 10.237294          | 4.479448           |
| std   | 2.275873           | 2.791840           | 2.461214           |
| min   | 1.040000           | 0.000000           | 0.000000           |
| 25%   | 7.520000           | 8.500000           | 3.000000           |
| 50%   | 9.050000           | 10.300000          | 4.000000           |
| 75%   | 10.590000          | 12.100000          | 6.000000           |
| max   | 17.770000          | 20.000000          | 20.000000          |

|       | total intl charge | customer service calls |
|-------|-------------------|------------------------|
| count | 3333.000000       | 3333.000000            |
| mean  | 2.764581          | 1.562856               |
| std   | 0.753773          | 1.315491               |
| min   | 0.000000          | 0.000000               |
| 25%   | 2.300000          | 1.000000               |
| 50%   | 2.780000          | 1.000000               |
| 75%   | 3.270000          | 2.000000               |
| max   | 5.400000          | 9.000000               |

-----  
Number of duplicates: 0

From the analysis of our dataset:

- There are 3333 rows and, 21 columns.
- There are no missing values.
- There are both categorical and numeric features.
- The numeric features are not all in the same scale.
- There are no duplicates in the dataset.

In [5]:

```
# Creating a dataframe to display datatypes and, the unique values.
desc = []
for i in df.columns:
    desc.append([
        i,
        df[i].dtypes,
        df[i].nunique(),
    ])

pd.DataFrame(data = desc, columns=['Feature', 'Dtypes', 'Sample_Unique'])
```

Out[5]:

|    | Feature                | Dtypes  | Sample_Unique |
|----|------------------------|---------|---------------|
| 0  | state                  | object  | 51            |
| 1  | account length         | int64   | 212           |
| 2  | area code              | int64   | 3             |
| 3  | phone number           | object  | 3333          |
| 4  | international plan     | object  | 2             |
| 5  | voice mail plan        | object  | 2             |
| 6  | number vmail messages  | int64   | 46            |
| 7  | total day minutes      | float64 | 1667          |
| 8  | total day calls        | int64   | 119           |
| 9  | total day charge       | float64 | 1667          |
| 10 | total eve minutes      | float64 | 1611          |
| 11 | total eve calls        | int64   | 123           |
| 12 | total eve charge       | float64 | 1440          |
| 13 | total night minutes    | float64 | 1591          |
| 14 | total night calls      | int64   | 120           |
| 15 | total night charge     | float64 | 933           |
| 16 | total intl minutes     | float64 | 162           |
| 17 | total intl calls       | int64   | 21            |
| 18 | total intl charge      | float64 | 162           |
| 19 | customer service calls | int64   | 10            |
| 20 | churn                  | bool    | 2             |


**Checking for Outliers.**



In [6]:

*# plotting a boxplot of each column to view outliers.*

```
numeric_columns = df.select_dtypes(include=['int64', 'float64'])
plt.figure(figsize=(20, 20))
plt.boxplot(numeric_columns.values, vert=False) # Show boxplots horizontally
plt.title("Box plot of all numeric columns in the dataset")
plt.xlabel("Value")
plt.yticks(range(1, len(numeric_columns.columns) + 1), numeric_columns.columns)
plt.show()
```



In [7]:

```
# Removing outliers
for col in numeric_columns:
    q1 = df[col].quantile(0.20)
    q3 = df[col].quantile(0.80)
    iqr = q3 - q1
    range_low = q1 - 1.5 * iqr
    range_high = q3 + 1.5 * iqr
    # Filtering the dataset in-place
    df = df.loc[(df[col] > range_low) & (df[col] < range_high)]

data = df
data.shape
```

Out[7]:

(3218, 21)

### 3. Univariate analysis


Exploring some of the variables in the data set looking for patterns of response to the variable.

In [8]:

```
# plot a bar graph to show the distribution of churn
data['churn'].value_counts().plot(kind='bar')
plt.xlabel('churn')
plt.ylabel('count')
plt.title('Distribution of Churn')
data['churn'].value_counts(normalize=True)
```

Out[8]:

```
False    0.85519
True     0.14481
Name: churn, dtype: float64
```



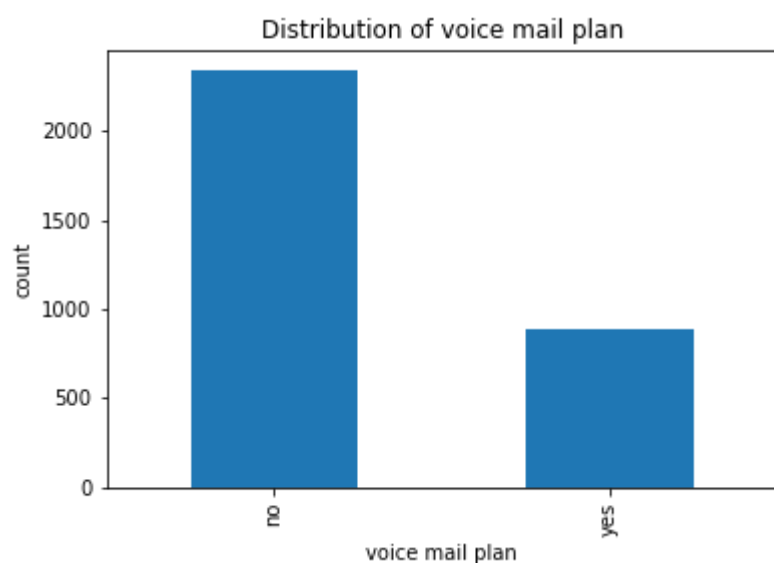
The bar graph above shows the distribution of churn where we see 14% of the customers churned.

In [9]:

```
# Distribution of voice mail plan  
  
data['voice mail plan'].value_counts().plot(kind='bar')  
plt.xlabel('voice mail plan')  
plt.ylabel('count')  
plt.title('Distribution of voice mail plan')  
data['voice mail plan'].value_counts(normalize=True)
```

Out[9]:

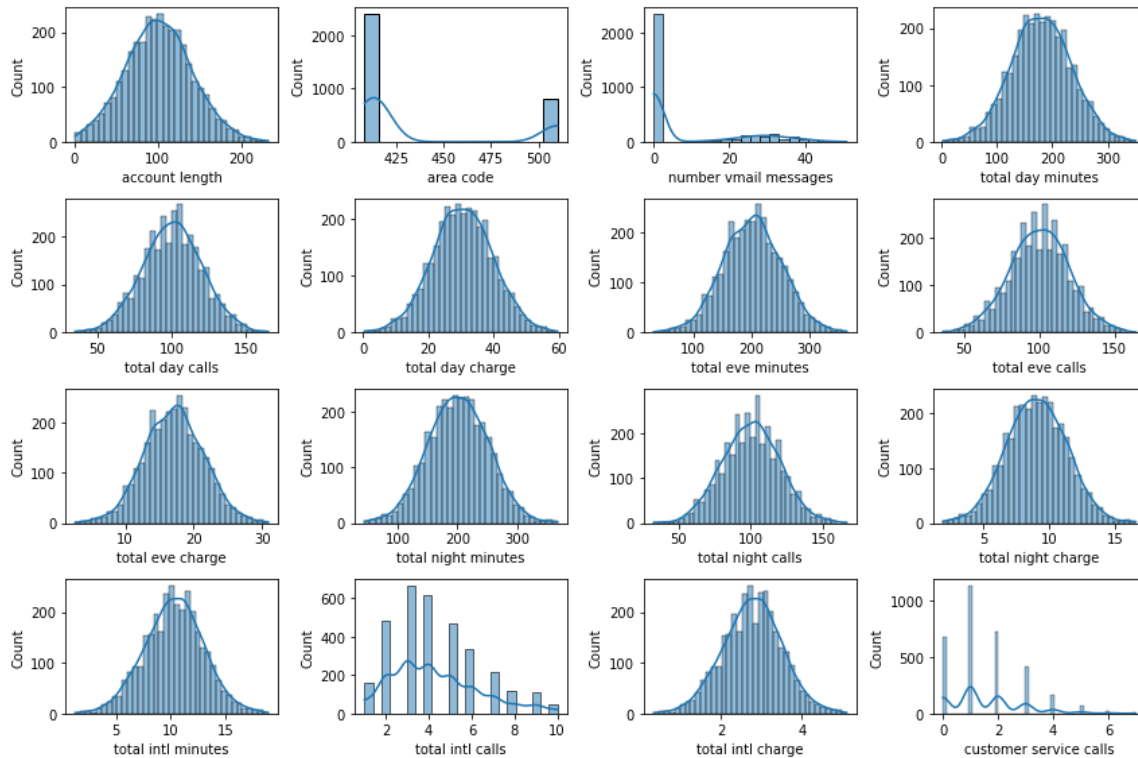
```
no      0.725606  
yes     0.274394  
Name: voice mail plan, dtype: float64
```



From the plot above we see that most customers (73%) did not have a voice mail plan.

In [10]:

```
# Checking the distribution of columns.
plt.figure(figsize=(12, 8))
for i, col in enumerate(numeric_columns, 1):
    plt.subplot(4, 4, i)
    sns.histplot(df[col], kde=True)
    plt.xlabel(col)
plt.tight_layout()
plt.show();
```



As seen from the distributions above most of the numeric features are normally distributed and the features will be good for modelling.

## 3.2 Bivariate analysis.

In this analysis, we will be examining the bivariate relationships between the features we have in relation to churn.

In [11]:

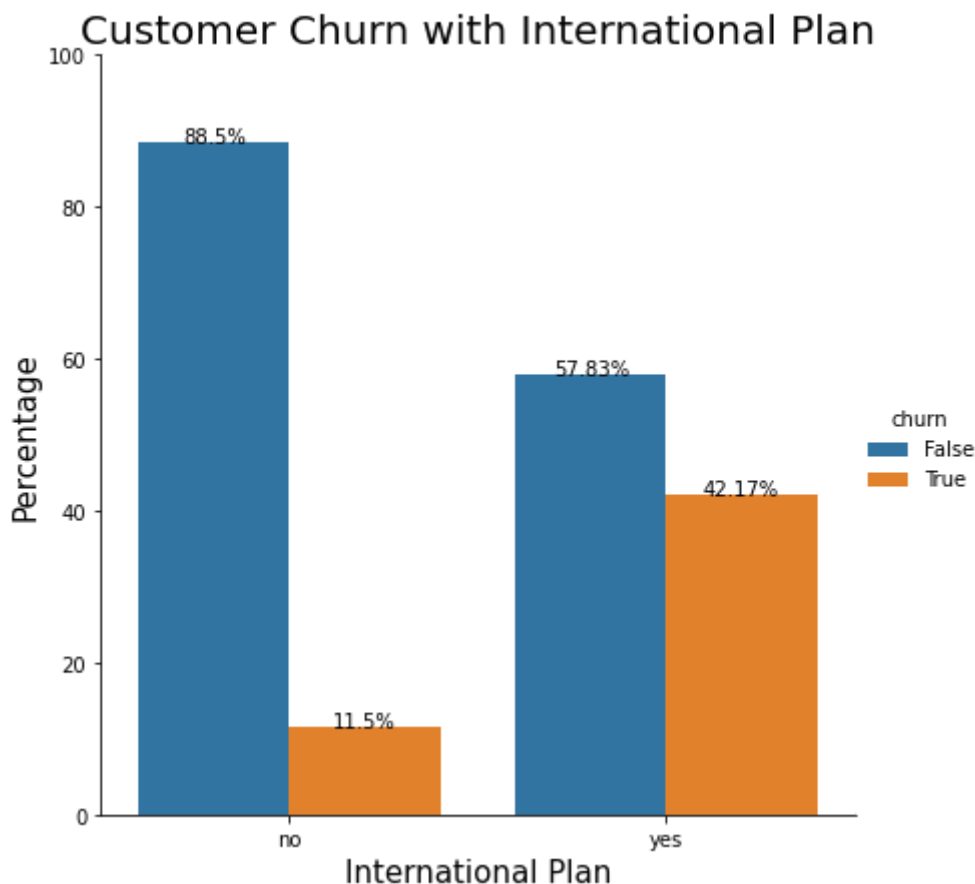
```
# A function that plots a feature in relation to churn.
def plot_churn_with_international_plan(data, x, y):
    df1 = data.groupby(x)['churn'].value_counts(normalize=True)
    df1 = df1.mul(100)
    df1 = df1.rename('Percentage').reset_index()

    ax = sns.catplot(x=x, y='Percentage', hue=y, kind='bar', data=df1, height=6)
    ax.set(ylim=(0, 100))

    for p in ax.ax.patches:
        txt = str(p.get_height().round(2)) + '%'
        txt_x = p.get_x() + p.get_width() / 2
        txt_y = p.get_height()
        ax.ax.text(txt_x, txt_y, txt, ha='center')
```

In [12]:


```
# A bar plot of international plan with churn
plot_churn_with_international_plan(data, 'international plan', 'churn')
plt.title('Customer Churn with International Plan', fontsize=20)
plt.xlabel('International Plan', fontsize=15)
plt.ylabel('Percentage', fontsize=15)
plt.show()
```



From the plot above the higher percentage of the customers who churned had an international plan.

In [13]:


```
# A bar plot of total day minute with churn
plot_churn_with_international_plan(data, 'customer service calls', 'churn')
plt.title('Customer Churn with customer service calls', fontsize=20)
plt.xlabel('Customer Service Calls', fontsize=15)
plt.ylabel('Percentage', fontsize=15)
plt.show()
```



From the plot above, a customer who had more customer service call are most likely to churn. This implies the costumers who did not have a customer service call are most likely satisfied from the services offered by the company.

In [14]:

```
# A bar plot of international calls with churn
plot_churn_with_international_plan(data, 'total intl calls', 'churn')
plt.title('Customer Churn with Total International call', fontsize=20)
plt.xlabel('International calls', fontsize=15)
plt.ylabel('Percentage', fontsize=15)
plt.show()
```

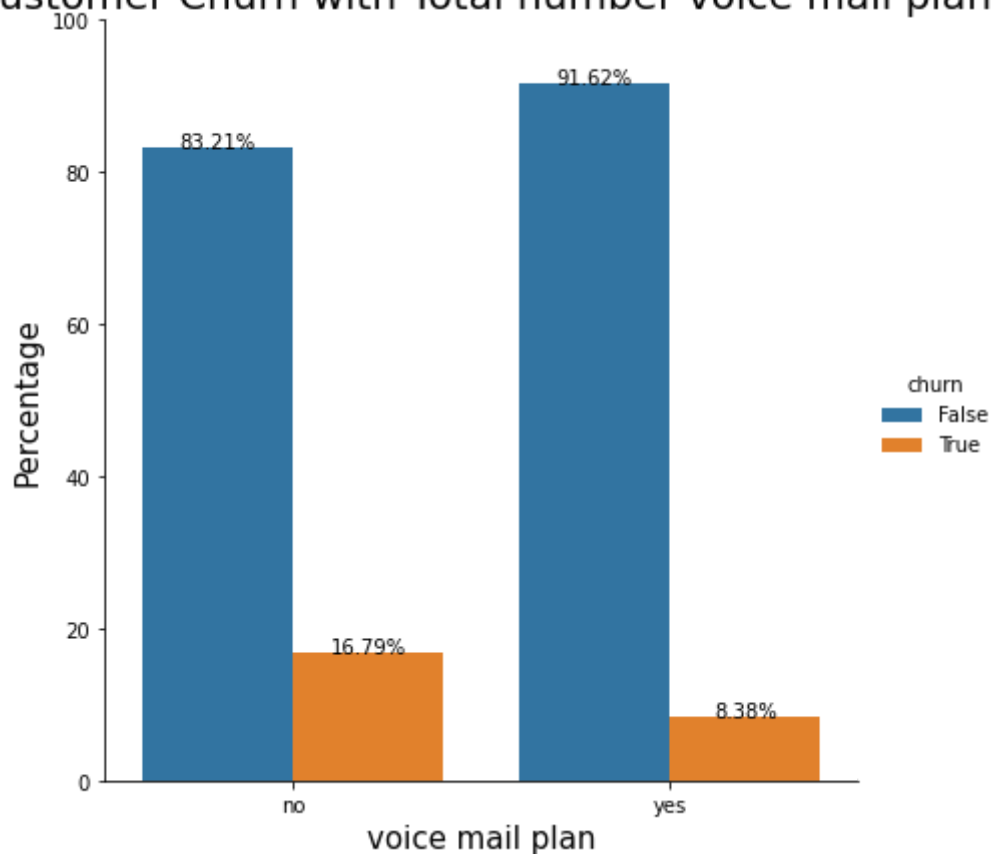


From the plot high number of international calls contributed to low churn rate.

In [15]:

```
# A bar plot of number vmail messages with churn
plot_churn_with_international_plan(data, 'voice mail plan', 'churn')
plt.title('Customer Churn with Total number voice mail plan', fontsize=20)
plt.xlabel('voice mail plan', fontsize=15)
plt.ylabel('Percentage', fontsize=15)
plt.show()
```

Customer Churn with Total number voice mail plan




From the plot above customers with a voice mail plan have a low churn rate.



In [16]:

```
# A scatter plot to show relationship between day minute and day calls with churn
sns.scatterplot(data=data, x='total day minutes', y='total day calls', hue='churn')
plt.xlabel('Total Day Minutes')
plt.ylabel('Total Day Calls')
plt.title('Relationship between Day Minutes, Day Calls, and Churn')
plt.show()
```



From the scatter plot above high total day minutes potentially leads to high churn rate.

## 4. Data Preprocessing

The steps for our data preprocessing include:

- Remove unnecessary columns ('Phone number')
- Check for correlation using heatmaps.
- Convert object binary columns into binary encoding of 0's and 1's.
- One-hot encode categorical columns.
- Store the target column, 'churn', in a separate variable and remove it from the dataframe
- Split the data into training and test sets.

we can first drop the features that will not be useful for the model, Phone number is a unique identifier of a customer and it is not required in our model.

In [17]:

```
# selecting the necessary columns for the model.
data = data.drop('phone number', axis=1)
```

In [18]:

```
# converting into binary columns.
data['international plan'] = df['international plan'].map({'yes':1, 'no':0})
data['voice mail plan'] = df['voice mail plan'].map({'yes':1, 'no':0})
data.head()
```

Out[18]:

|   | state | account length | area code | international plan | voice mail plan | number vmail messages | total day minutes | total day calls | total day charge | total eve minutes | total eve charge |
|---|-------|----------------|-----------|--------------------|-----------------|-----------------------|-------------------|-----------------|------------------|-------------------|------------------|
| 0 | KS    | 128            | 415       | 0                  | 1               | 25                    | 265.1             | 110             | 45.07            | 197.4             | 9.91             |
| 1 | OH    | 107            | 415       | 0                  | 1               | 26                    | 161.6             | 123             | 27.47            | 195.5             | 10.44            |
| 2 | NJ    | 137            | 415       | 0                  | 0               | 0                     | 243.4             | 114             | 41.38            | 121.2             | 1.87             |
| 3 | OH    | 84             | 408       | 1                  | 0               | 0                     | 299.4             | 71              | 50.90            | 61.9              | 8.26             |
| 4 | OK    | 75             | 415       | 1                  | 0               | 0                     | 166.7             | 113             | 28.34            | 148.3             | 1.87             |


In [19]:

```
# Assign the 'churn' column to labels
labels = data.churn

# Drop the 'churn' column from the dataframe.
predictors = data.drop('churn',axis=1)
```

In [20]:

```
# a heatmap that shows correlation.
plt.figure(figsize=(20, 20))
cmap = sns.diverging_palette(220,20,n=200)
heatmap = sns.heatmap(
    predictors.corr(),vmin=-1, vmax=1,center = 0,
    annot=True,cmap=cmap,linewidths=2, linecolor='black')
heatmap.set_title('Correlation Heatmap', fontdict={'fontsize':18}, pad=12);
```



In [21]:

```
# correlations with churn
data.corr().churn.sort_values(ascending = False)
```

Out[21]:

```
churn                1.000000
international plan    0.258294
total day minutes     0.206320
total day charge      0.206318
customer service calls 0.201696
total eve minutes     0.093162
total eve charge      0.093151
total intl charge     0.056676
total intl minutes    0.056659
total night charge    0.037257
total night minutes   0.037243
account length        0.016708
total day calls        0.015308
area code             0.010032
total night calls     0.007787
total eve calls        0.004030
total intl calls      -0.070657
number vmail messages -0.094225
voice mail plan       -0.106604
Name: churn, dtype: float64
```

Here we see that all minutes and charge features are perfectly correlated ( $r = 1$ ). This implies charge is usually based on minutes.

In [22]:

```
# importing OneHotEncoder
from sklearn.preprocessing import OneHotEncoder

# One-hot encode categorical columns.
ohe = OneHotEncoder(sparse = False, handle_unknown = "ignore")

# fit ohe on small train data
ohe.fit(predictors[['state']])

# access the column names of the states
col_names = ohe.categories_[0]

# make a df with encoded states
train_state_encoded = pd.DataFrame(ohe.transform(predictors[["state"]]),
                                   index = predictors.index,
                                   columns = col_names)

# combine encoded features
predictors = pd.concat([predictors.drop("state", axis = 1), train_state_encoded],
```

In [23]:

```
# Split the data into training and test sets
```

```
X_train, X_test, y_train, y_test = train_test_split(predictors, labels, test_size :
```

In [24]:

```
X_train.head()
```

Out[24]:

|             | account<br>length | area<br>code | international<br>plan | voice<br>mail<br>plan | number<br>vmail<br>messages | total<br>day<br>minutes | total<br>day<br>calls | total<br>day<br>charge | total<br>eve<br>minutes | total<br>eve<br>calls |
|-------------|-------------------|--------------|-----------------------|-----------------------|-----------------------------|-------------------------|-----------------------|------------------------|-------------------------|-----------------------|
| <b>2401</b> | 126               | 415          | 1                     | 0                     | 0                           | 239.7                   | 87                    | 40.75                  | 281.7                   | 92                    |
| <b>991</b>  | 50                | 415          | 0                     | 1                     | 35                          | 192.6                   | 97                    | 32.74                  | 135.2                   | 101                   |
| <b>2725</b> | 51                | 408          | 0                     | 0                     | 0                           | 169.3                   | 111                   | 28.78                  | 139.5                   | 69                    |
| <b>778</b>  | 115               | 415          | 0                     | 1                     | 26                          | 170.5                   | 107                   | 28.99                  | 217.2                   | 77                    |
| <b>235</b>  | 139               | 510          | 0                     | 0                     | 0                           | 134.4                   | 106                   | 22.85                  | 211.3                   | 98                    |

5 rows × 69 columns

since now all the columns are numeric we can proceed and create a logistic baseline model

## 5. Modeling

At this step I will build three models to and evaluate their performance. I will then choose the model that has the best performance and tune it to aiming to get better performance on the predictions made.

The models that I will try out are:

- 1.Logistic Regression
- 2.Decision Trees
- 3.Random Forest
- 4.Final Tuned model.

### 5.1 Logistic Regression


From the analysis it was evident that there was a class imbalance. Because of class the imbalance, we should add some kind of resampling step. Specifically we'll use SMOTE from imblearn.

In [25]:

```
# checking for imbalance
data['churn'].value_counts().plot(kind='bar')
plt.xlabel('churn')
plt.ylabel('count')
plt.title('Distribution of Churn')
data['churn'].value_counts(normalize=True)
```

Out[25]:

```
False    0.85519
True     0.14481
Name: churn, dtype: float64
```



Since the features are on different scale we can first scale and then fit then use a resampling method on the class imbalance.

In [26]:

```
# Scale X_train and X_test using StandardScaler
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Ensure X_train and X_test are scaled DataFrames
X_train1 = pd.DataFrame(X_train_scaled, columns=X_train.columns)
X_test1 = pd.DataFrame(X_test_scaled, columns=X_train.columns)

# Fit SMOTE to training data
X_train_resampled, y_train_resampled = SMOTE().fit_resample(X_train_scaled, y_train)
```

In [27]:

```
# Fit a model
logreg = LogisticRegression(fit_intercept=False, solver='liblinear')
model_log = logreg.fit(X_train_resampled, y_train_resampled)
```

**Evaluation.**

For evaluation we use Recall as it measures the proportion of correctly predicted positive instances out of all actual positive instances. Recall focuses on the model's ability to correctly identify positive instances and is useful when the cost of false negatives is high.

In [28]:

```
# Evaluating Logistic regression model.
print(f'Train Accuracy: {accuracy_score(y_train, model_log.predict(X_train1))}')
print(f'Test Accuracy: {accuracy_score(y_test, model_log.predict(X_test1))}')
print('-----')
print(f'Train Recall: {recall_score(y_train, model_log.predict(X_train1))}')
print(f'Test Recall: {recall_score(y_test, model_log.predict(X_test1))}')
```

```
Train Accuracy: 0.641110650642354
Test Accuracy: 0.6285714285714286
```

```
-----
Train Recall: 0.8854748603351955
Test Recall: 0.8055555555555556
```

From the above evaluation metrics the accuracy score is not satisfactory but the recall score is fairly much better. We can try to use another model, see if it improves in metrics.

## 5.2 Decision Trees

In [29]:

```
# Instantiate and fit a DecisionTreeClassifier
tree_clf = DecisionTreeClassifier(criterion='gini', max_depth = 5, random_state=42)
tree_clf.fit(X_train,y_train)
```

Out[29]:

```
DecisionTreeClassifier(max_depth=5, random_state=42)
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

In [30]:

```
# Evaluation of model accuracy and recall.
print(f'Train Accuracy: {accuracy_score(y_train, tree_clf.predict(X_train))}')
print(f'Test Accuracy: {accuracy_score(y_test, tree_clf.predict(X_test))}')
print('-----')

print(f'Train Recall: {recall_score(y_train, tree_clf.predict(X_train))}')
print(f'Test Recall: {recall_score(y_test, tree_clf.predict(X_test))}')
```

```
Train Accuracy: 0.9481972648155823
Test Accuracy: 0.9341614906832298
```

```
-----
Train Recall: 0.7458100558659218
Test Recall: 0.7222222222222222
```

From the evaluation of the Decision tree the accuracy score is very good and Recall is fairly good we can fit yet another model to improve recall score.

## **Feature importance.**


we can examine to see which features were important in our decision tree model



In [31]:

```
# Defining a function that plots the most important features.
def plot_feature_importances(model):
    n_features = X_train.shape[1]
    sorted_idx = np.argsort(model.feature_importances_)[::-1] # Sort feature imp
    plt.figure(figsize=(15, 20))
    plt.barh(range(n_features), model.feature_importances_[sorted_idx], align='ce
    plt.yticks(np.arange(n_features), X_train.columns.values[sorted_idx])
    plt.xlabel('Feature Importance')
    plt.ylabel('Feature')
    plt.title('Feature Importances')
    plt.show()

plot_feature_importances(tree_clf)
```



Here as we can see from our model the 5 most important features are:

1. Total day minutes.

2. Customer service calls.
3. Total day charge.
4. Total evening charge.
5. Total international calls.

## 5.3 Random Forest.

In [32]:

```
# Instantiate and fit a RandomForestClassifier
forest = RandomForestClassifier(n_estimators=100, max_depth=5, random_state=42)
forest.fit(X_train,y_train)
```

Out[32]:

```
RandomForestClassifier(max_depth=5, random_state=42)
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

In [33]:

```
# Evaluation.
print(f'Train Accuracy: {accuracy_score(y_train, forest.predict(X_train))}')
print(f'Test Accuracy: {accuracy_score(y_test, forest.predict(X_test))}')
print('-----')


print(f'Train Recall: {recall_score(y_train, forest.predict(X_train))}')
print(f'Test Recall: {recall_score(y_test, forest.predict(X_test))}')
```

```
Train Accuracy: 0.886033982594281
Test Accuracy: 0.8832298136645963
-----
Train Recall: 0.23184357541899442
Test Recall: 0.12962962962962962
```

The model performed very poorly on recall and the accuracy is a bit low as compared to decision tree, Tuning this model can potentially improve the metrics. so we are going to try hyperparameter tuning using grid search.

In [34]:

```
plot_feature_importances(forest)
```



from our Random forest model the 5 most important features are:

- 1. Total day minutes.
- 2. Total day charge.
- 3. Customer service calls.
- 4. Internatinal plan
- 5. Total evening charge.

### Hyperparameter Tuning

we can perform a grid search on the forest to identify the best parameters to use for our model

In [35]:

```
# Defining a parameter grid
dt_param_grid = {
    "criterion": ["gini", "entropy"],
    "max_depth": [None, 2, 3, 4, 5, 6],
    "min_samples_split": [2, 5, 10],
    "min_samples_leaf": [1, 2, 3, 4, 5, 6],
}
```

In [36]:

```
from sklearn.model_selection import GridSearchCV
# Instantiate GridSearchCV
dt_grid_search = GridSearchCV(tree_clf, dt_param_grid, cv=3, return_train_score=T
# Fit to the data
dt_grid_search.fit(X_train, y_train)
```

Out[36]:

```
GridSearchCV(cv=3,
             estimator=DecisionTreeClassifier(max_depth=5, random_state=42),
             param_grid={'criterion': ['gini', 'entropy'],
                          'max_depth': [None, 2, 3, 4, 5, 6],
                          'min_samples_leaf': [1, 2, 3, 4, 5, 6],
                          'min_samples_split': [2, 5, 10]},
             return_train_score=True)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [37]:

```
# Mean training score
dt_gs_training_score = np.mean(dt_grid_search.cv_results_["mean_train_score"])

# Mean test score
dt_gs_testing_score = dt_grid_search.score(X_test, y_test)

print(f"Mean Training Score: {dt_gs_training_score :.2%}")
print(f"Mean Test Score: {dt_gs_testing_score :.2%}")
print("Best Parameter Combination Found During Grid Search:")
print('-----')

print(f'Train Recall: {recall_score(y_train, dt_grid_search.predict(X_train))}')
print(f'Test Recall: {recall_score(y_test, dt_grid_search.predict(X_test))}')

dt_grid_search.best_params_
```

Mean Training Score: 93.54%

Mean Test Score: 95.03%

Best Parameter Combination Found During Grid Search:

-----

Train Recall: 0.7849162011173184

Test Recall: 0.7222222222222222

Out[37]:

```
{'criterion': 'entropy',
 'max_depth': 6,
 'min_samples_leaf': 3,
 'min_samples_split': 10}
```

## Model Summary

I used three classification models to predict churn and a more tuned model to improve the evaluation metrics i.e:

- Logistic regression which had a: -Train Accuracy: 65% -Test Accuracy: 63% -Train Recall: 89% -Test Recall: 80%
- Decision Tree which had a: -Train Accuracy: 95% -Test Accuracy: 93% -Train Recall: 75% -Test Recall: 72%
- Random Forest with a: -Train Accuracy: 89% -Test Accuracy: 88% -Train Recall: 23% -Test Recall: 12%
- Tuned Random Forest with a: -Train Accuracy: 94% -Test Accuracy: 95% -Train Recall: 78% -Test Recall: 72% The best parameters for random forest are criterion 'entropy', Maximum depth of 6, minimum sample leaf of 3 and a minimum sample split of 10

## Conclusion

The final model that will be used to predict customer churn is Random Forest with the tuned hyperparameters (criterion 'entropy', Maximum depth of 6, minimum sample leaf of 3 and a minimum sample split of 10) as it has the least number of false negatives with a better recall score. The most important features that lead to customer churn are:

- Total day minutes.
- Total day charge.
- Customer service calls.

- International plan.
- Total evening charge.

## Recommendations.

From the findings

- A customer who had more customer service call are most likely to churn. This implies the costumers who did not have a customer service call are most likely satisfied from the services offered by the company.
- A higher percentage of the customers who churned had an international plan.
- High total day minutes spent potentially leads to high churn rate, this implies that a customer was charged based on the total number of minutes spent.
- Total evening charge and night minutes also increased the churn rate. Therefore from the findings I would reccomend SyriaTel to:

1. Improve the quality of customer service provided
2. Review the international plan.
3. Review on the prices how they charge customers based on the total number of minutes.

## Limitations.

- The dataset used has limited number of features while there could be other factors that contribute to customer churn.
- The classification models used in this analysis make certain assumptions, such as independence, and normality. Violations of these assumptions could impact the accuracy and reliability of the results.

In [ ]: